

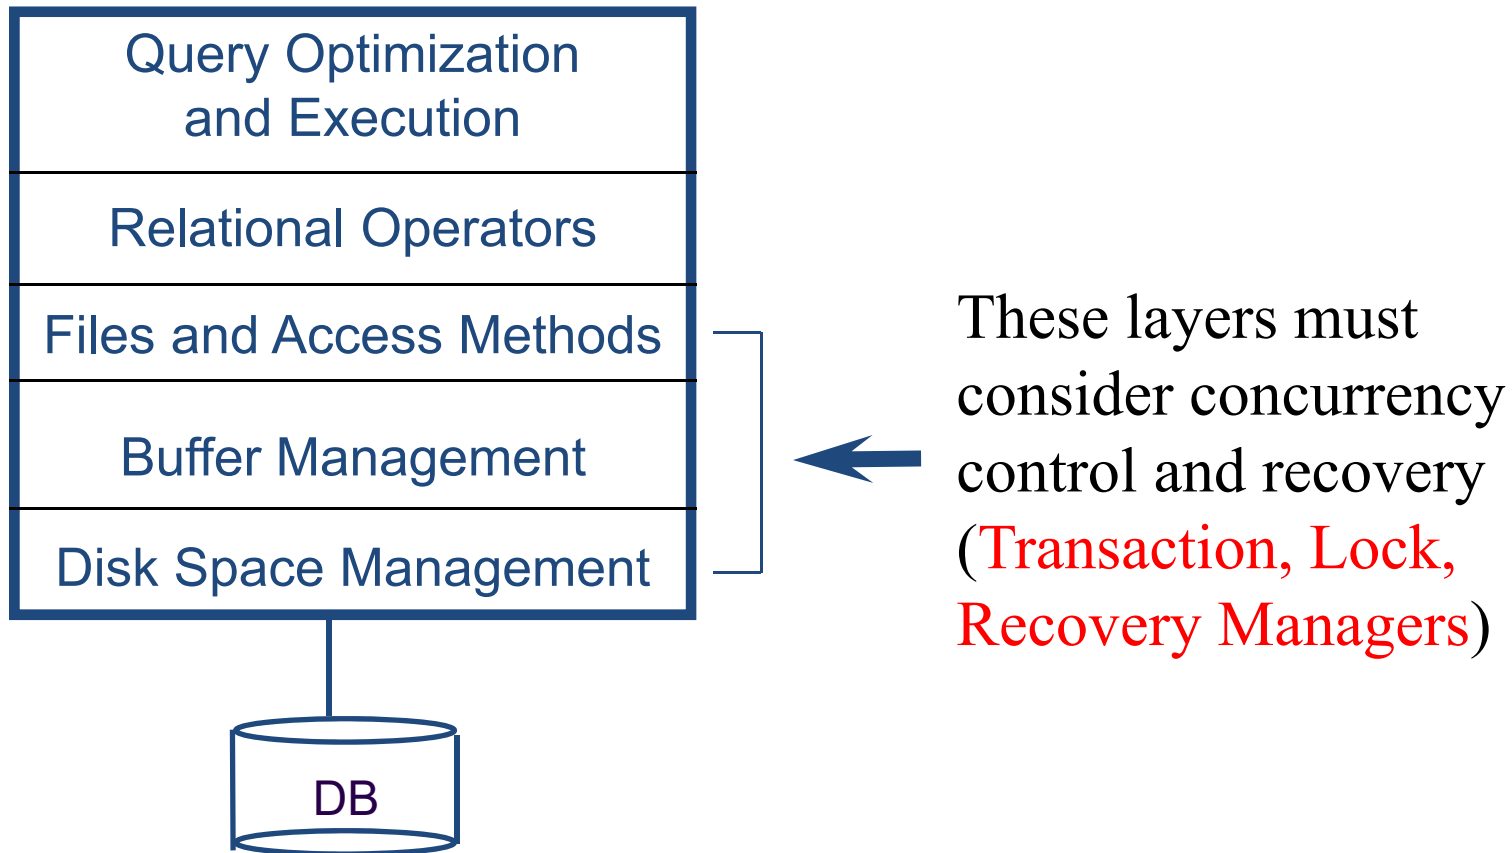
Database Management Systems

Sumayyea Salahuddin (Lecturer)
Dept. of Computer Systems Eng.
UET Peshawar

Overview

- Database **transactions** and their properties
- What **concurrency control** is and what role it plays in maintaining the database's integrity
- What **locking** methods are and how they work
- How **stamping** methods are used for concurrency control
- How **optimistic methods** are used for concurrency control
- How database **recovery management** is used to maintain database integrity

DBMS Structure



What is Transaction?

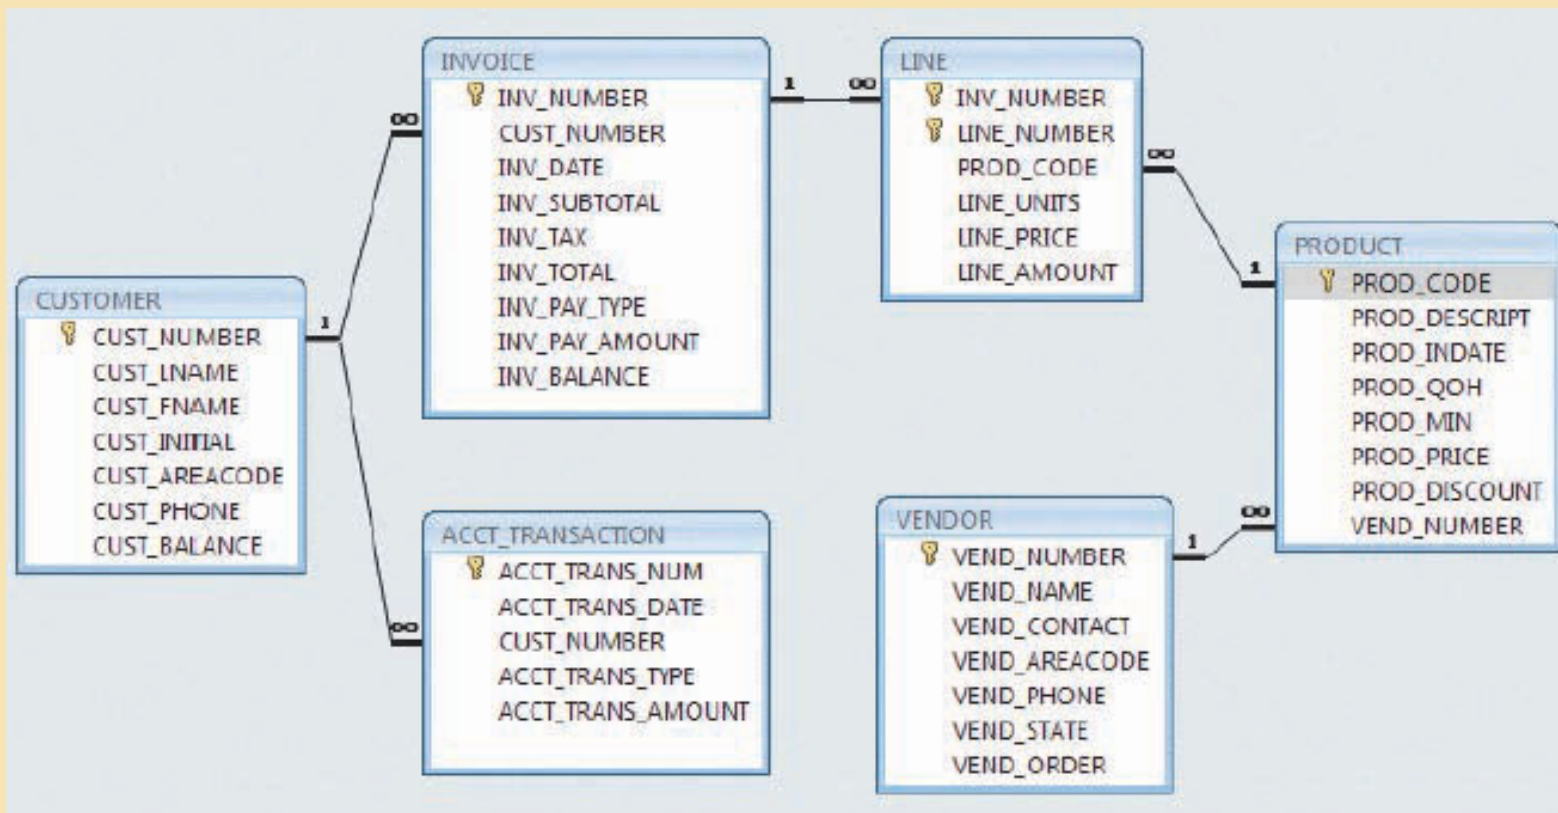
- Logical unit of work that must be either entirely completed or aborted
- Successful transaction changes database from one consistent state to another
 - One in which all data integrity constraints are satisfied
- Most real-world database transactions are formed by two or more database requests
 - Equivalent of a single SQL statement in an application program or transaction
 - if a transaction is composed of two UPDATE statements and one INSERT statement, the transaction uses three database requests. In turn, each database request generates several input/output (I/O) operations that read from or write to physical storage media.

What is Transaction? (Cont...)

- A transaction is any action that reads from and/or writes to a database & may consist of:
 - Simple SELECT statement to generate a list of table contents
 - Series of related UPDATE statements to change values of attributes in various tables
 - Series of INSERT statements to add rows to one or more tables
 - Combination of SELECT, UPDATE, and INSERT statements
- Example in next slide

Sample Relational Database

FIGURE 10.1 The Ch10_SaleCo database relational diagram



Example: Sales Transaction

- START TRANSACTION;
- INSERT INTO INVOICE VALUES (1009, 10016, '18-Jan-2008', 256.99, 20.56, 277.55, 'cred', 0.00, 277.55);
- INSERT INTO LINE VALUES (1009, 1, '89-WRE-Q', 1, 256.99, 256.99);
- UPDATE PRODUCT SET PROD_QOH = PROD_QOH - 1 WHERE PROD_CODE = '89-WRE-Q';
- UPDATE CUSTOMER SET CUST_BALANCE = CUST_BALANCE + 277.55 WHERE CUST_NUMBER = 10016;
- INSERT INTO ACCT_TRANSACTION VALUES (10007, '18-Jan-08', 10016, 'charge', 277.55);
- COMMIT/ROLLBACK;

Transaction Results

Table name: INVOICE

| INV_NUMBER | CUST_NUMBER | INV_DATE | INV_SUBTOTAL | INV_TAX | INV_TOTAL | INV_PAY_TYPE | INV_PAY_AMOUNT | INV_BALANCE |
|------------|-------------|-----------|--------------|---------|-----------|--------------|----------------|-------------|
| 1001 | 10014 | 16-Jan-08 | 54.02 | 4.30 | 59.31 | cc | 59.31 | 0.00 |
| 1002 | 10011 | 16-Jan-08 | 9.98 | 0.80 | 10.78 | cash | 10.78 | 0.00 |
| 1003 | 10012 | 16-Jan-08 | 270.70 | 21.00 | 292.30 | cc | 292.30 | 0.00 |
| 1004 | 10011 | 17-Jan-08 | 34.87 | 2.79 | 37.66 | cc | 37.66 | 0.00 |
| 1005 | 10010 | 17-Jan-08 | 70.44 | 5.04 | 75.00 | cc | 75.00 | 0.00 |
| 1006 | 10014 | 17-Jan-08 | 397.83 | 31.83 | 429.66 | cred | 100.00 | 329.66 |
| 1007 | 10015 | 17-Jan-08 | 34.97 | 2.80 | 37.77 | chil | 37.77 | 0.00 |
| 1008 | 10011 | 17-Jan-08 | 1033.08 | 82.65 | 1115.73 | cred | 500.00 | 615.73 |
| 1009 | 10018 | 18-Jan-08 | 258.99 | 20.56 | 277.55 | cred | 0.00 | 277.55 |

Table name: LINE

| INV_NUMBER | LINE_NUMBER | PROD_CODE | LINE_UNITS | LINE_PRICE | LINE_AMOUNT |
|------------|-------------|-----------|------------|------------|-------------|
| 1001 | 1 | 13-Q2P2 | 3 | 14.00 | 44.07 |
| 1001 | 2 | 23109-HB | 1 | 9.95 | 9.95 |
| 1002 | 1 | 54778-2T | 2 | 4.00 | 8.00 |
| 1003 | 1 | 233806D | 4 | 38.95 | 155.80 |
| 1003 | 2 | 1346-Q02 | 1 | 39.95 | 39.95 |
| 1003 | 3 | 13-Q2P2 | 6 | 14.00 | 84.00 |
| 1004 | 1 | 54778-2T | 3 | 4.99 | 14.97 |
| 1004 | 2 | 23109-HB | 2 | 9.95 | 19.90 |
| 1005 | 1 | PVC33DRT | 12 | 5.67 | 68.04 |
| 1005 | 1 | SM-18277 | 3 | 6.00 | 18.00 |
| 1005 | 2 | 2232QTY | 1 | 109.92 | 109.92 |
| 1005 | 3 | 23100-HB | 1 | 0.05 | 0.05 |
| 1005 | 4 | 89-WRE-Q | 1 | 256.99 | 256.99 |
| 1007 | 1 | 13-Q2P2 | 2 | 14.99 | 29.98 |
| 1007 | 2 | 54778-2T | 1 | 4.99 | 4.99 |
| 1008 | 1 | PVC33DRT | 5 | 5.07 | 25.35 |
| 1008 | 2 | WR31T13 | 4 | 119.85 | 479.40 |
| 1008 | 3 | 23109-HB | 1 | 9.95 | 9.95 |
| 1008 | 4 | 89-WRE-Q | 2 | 256.00 | 512.00 |
| 1009 | 1 | 89-WRE-Q | 1 | 256.99 | 256.99 |

Table name: PRODUCT

| PROD_CODE | PROD_DESCRIPTOR | PROD_INDATE | PROD_QOH | PROD_MIN | PROD_PRICE | PROD_DISCOUNT | VEND_NUMBER |
|-----------|----------------------------------|-------------|----------|----------|------------|---------------|-------------|
| 13-Q2P2 | Power painter, 15 psi, 3-nozzle | 03-Nov-07 | 8 | 5 | 100.00 | 0.00 | 25505 |
| 13-Q2P2 | 7.25-in. pwr. saw blade | 13-Dec-07 | 32 | 15 | 14.99 | 0.05 | 21344 |
| 14-Q1AL3 | 9.00-in. pwr. saw blade | 13-Nov-07 | 10 | 12 | 17.49 | 0.00 | 21344 |
| 1516-Q02 | Hrd. cloth, 14 in., 2x50 | 15-Jan-08 | 15 | 8 | 39.95 | 0.00 | 23119 |
| 1550-Q01 | Hrd. cloth, 12-in., 3x50 | 15-Jan-08 | 20 | 5 | 43.99 | 0.00 | 23119 |
| 2232QTY | B&D jigsaw, 12 in. blade | 30-Dec-07 | 8 | 5 | 109.92 | 0.05 | 24288 |
| 2232QMF | B&D jigsaw, 8-in. blade | 24-Dec-07 | 6 | 5 | 99.87 | 0.05 | 24288 |
| 2238GPD | B&D cordless drill, 1/2-in. | 20-Jan-08 | 12 | 5 | 38.95 | 0.05 | 25505 |
| 23109-HB | Crew hammer | 20-Jan-08 | 23 | 10 | 9.95 | 0.10 | 21225 |
| 23114-AA | Sledge hammer, 12 lb. | 02-Jan-08 | 8 | 5 | 14.40 | 0.05 | |
| 54778-2T | Rat-tail file, 1/8-in. fine | 15-Dec-07 | 43 | 20 | 4.99 | 0.00 | 21344 |
| 89-WRE-Q | 1/2-in. chain saw, 10 in. | 07-Jan-08 | 11 | 5 | 256.99 | 0.05 | 24200 |
| PVC33DRT | PVC pipe, 3.5 in., 8 ft | 06-Jan-08 | 188 | 75 | 5.67 | 0.00 | |
| SM-18277 | 1.25-in. metal screw, 25 | 01-Mar-08 | 172 | 75 | 6.99 | 0.00 | 21225 |
| SW-23116 | 2.5-in. wd. screw, 50 | 24-Feb-08 | 237 | 100 | 8.45 | 0.00 | 21231 |
| WR31T13 | Steel netting, 4x8x1.6", 5" mesh | 17-Jan-08 | 18 | 5 | 119.95 | 0.10 | 25595 |

Table name: CUSTOMER

| CUST_NUMBER | CUST_LNAME | CUST_FNAME | CUST_INITIAL | CUST_AREACODE | CUST_PHONE | CUST_BALANCE |
|-------------|------------|------------|--------------|---------------|------------|--------------|
| 10010 | Romero | Alfred | A | 815 | 844-2573 | 0.00 |
| 10011 | Dunne | Leona | K | 713 | 894-1298 | 615.73 |
| 10012 | Smith | Kathy | W | 916 | 904-2296 | 0.00 |
| 10013 | Okrowski | Paul | F | 815 | 894-2180 | 0.00 |
| 10014 | Orlando | Myron | | 815 | 222-1872 | 0.00 |
| 10015 | O'Brien | Amy | B | 713 | 442-3381 | 0.00 |
| 10016 | Brown | James | Q | 815 | 297-1228 | 277.55 |
| 10017 | Williams | George | | 815 | 290-2558 | 0.00 |
| 10018 | Farriss | Anne | Q | 713 | 382-7185 | 0.00 |
| 10019 | Smith | Olette | K | 815 | 297-3809 | 0.00 |

Table name: ACCT TRANSACTION

| ACCT_TRANS_NUM | ACCT_TRANS_DATE | CUST_NUMBER | ACCT_TRANS_TYPE | ACCT_TRANS_AMOUNT |
|----------------|-----------------|-------------|-----------------|-------------------|
| 10003 | 17-Jan-08 | 10014 | charge | 329.66 |
| 10004 | 17-Jan-08 | 10011 | charge | 615.73 |
| 10005 | 29-Jan-08 | 10014 | payment | 329.66 |
| 10007 | 18-Jan-08 | 10018 | charge | 277.55 |

Evaluating Transaction Results

- Not all transactions update database
- SQL code represents a transaction because database was accessed
- Improper or incomplete transactions can have devastating effect on database integrity
 - Some DBMSs provide means by which user can define enforceable constraints
 - Other integrity rules are enforced automatically by the DBMS
 - Primary key integrity
 - Referential integrity
 - Entity integrity
- No semantic/logical checking
- Transactions that violate integrity constraints are aborted

Transaction Properties

- **Atomicity**
 - All operations of a transaction must be completed
- **Consistency**
 - Permanence of database's consistent state
- **Isolation**
 - Data used during transaction cannot be used by second transaction until the first is completed
- **Durability**
 - Once transactions are committed, they cannot be undone
- **Serializability**
 - Concurrent execution of several transactions yields consistent results

Multuser databases subject to multiple concurrent transactions

Transaction Management with SQL

- ANSI has defined standards that govern SQL database transactions
- Transaction support is provided by two SQL statements: **COMMIT** and **ROLLBACK**
- Transaction sequence must continue until:
 - COMMIT statement is reached
 - ROLLBACK statement is reached
 - End of program is reached (Equivalent to COMMIT)
 - Program is abnormally terminated (Equivalent to ROLLBACK)

Transaction & Concurrent Execution

- **Transaction Manager** controls the execution of transactions.
- User program may carry out many operations on the data retrieved from the database, but the DBMS is only concerned about what data is read/written from/to the database.
- Concurrent execution of multiple transactions is essential for good performance.
 - Disk is the bottleneck (slow, frequently used)
 - Must keep CPU busy w/many queries
 - Better response time

The Transaction Log


- Keeps track of all transactions that **update the database**
- Useful for recovering database
- **Transaction log** stores:
 - A record for the beginning of transaction
 - For each transaction component:
 - Type of operation being performed (update, delete, insert)
 - Names of objects affected by transaction
 - "Before" and "after" values for updated fields
 - Pointers to previous and next transaction log entries for the same transaction
 - Ending (COMMIT) of the transaction

Sample Transaction Log

TABLE
10.1

A Transaction Log

| TRL_ID | TRX_NUM | PREV_PTR | NEXT_PTR | OPERATION | TABLE | ROW ID | ATTRIBUTE | BEFORE VALUE | AFTER VALUE |
|--------|---------|----------|----------|-----------|-------------------------|----------|--------------|--------------|-------------|
| 341 | 101 | Null | 352 | START | ****Start Transaction | | | | |
| 352 | 101 | 341 | 363 | UPDATE | PRODUCT | 1558-QW1 | PROD_QOH | 25 | 23 |
| 363 | 101 | 352 | 365 | UPDATE | CUSTOMER | 10011 | CUST_BALANCE | 525.75 | 615.73 |
| 365 | 101 | 363 | Null | COMMIT | **** End of Transaction | | | | |



TRL_ID = Transaction log record ID **PTR** = Pointer to a transaction log record ID
TRX_NUM = Transaction number
 (Note: The transaction number is automatically assigned by the DBMS.)

```

START TRANSACTION;
UPDATE PRODUCT SET PROD_QOH = 23 WHERE PROD_ID = '1558-QW1';
UPDATE CUSTOMER SET CUST_BALANCE = 615.73 WHERE CUST_ID =
10011;
COMMIT;
  
```

Concurrency Control

- Coordination of simultaneous transaction execution in a multiprocessing database
- Objective is to ensure serializability of transactions in a multiuser environment
- Simultaneous execution of transactions over a shared database can create several data integrity and consistency problems
 - Lost updates
 - Uncommitted data
 - Inconsistent retrievals

Lost Updates

- **Lost update** problem:
 - Two concurrent transactions update same data element
 - One of the update is lost
 - Overwritten by the other transaction

TABLE
10.2

Two Concurrent Transactions to Update QOH

| TRANSACTION | COMPUTATION |
|------------------------|---|
| T1: Purchase 100 units | $\text{PROD_QOH} = \text{PROD_QOH} + 100$ |
| T2: Sell 30 units | $\text{PROD_QOH} = \text{PROD_QOH} - 30$ |

Lost Updates (Cont...)

TABLE
10.3

Serial Execution of Two Transactions

| TIME | TRANSACTION | STEP | STORED VALUE |
|------|-------------|-------------------------------|--------------|
| 1 | T1 | Read PROD_QOH | 35 |
| 2 | T1 | $\text{PROD_QOH} = 35 + 100$ | |
| 3 | T1 | Write PROD_QOH | 135 |
| 4 | T2 | Read PROD_QOH | 135 |
| 5 | T2 | $\text{PROD_QOH} = 135 - 30$ | |
| 6 | T2 | Write PROD_QOH | 105 |

TABLE
10.4

Lost Updates

| TIME | TRANSACTION | STEP | STORED VALUE |
|------|-------------|-------------------------------|--------------|
| 1 | T1 | Read PROD_QOH | 35 |
| 2 | T2 | Read PROD_QOH | 35 |
| 3 | T1 | $\text{PROD_QOH} = 35 + 100$ | |
| 4 | T2 | $\text{PROD_QOH} = 35 - 30$ | |
| 5 | T1 | Write PROD_QOH (Lost update) | 135 |
| 6 | T2 | Write PROD_QOH | 5 |

Uncommitted Data

- Uncommitted data phenomenon:
 - Two transaction executed concurrently
 - First transaction rolled back after second already accessed uncommitted data

TABLE
10.5

Transactions Creating Uncommitted Data Problem

| TRANSACTION | COMPUTATION |
|------------------------|---|
| T1: Purchase 100 units | PROD_QOH = PROD_QOH + 100 (Rolled back) |
| T2: Sell 30 units | PROD_QOH = PROD_QOH 30 |

Uncommitted Data (Cont...)

TABLE
10.6

Correct Execution of Two Transactions

| TIME | TRANSACTION | STEP | STORED VALUE |
|------|-------------|-------------------------------|--------------|
| 1 | T1 | Read PROD_QOH | 35 |
| 2 | T1 | $\text{PROD_QOH} = 35 + 100$ | |
| 3 | T1 | Write PROD_QOH | 135 |
| 4 | T1 | *****ROLLBACK***** | 35 |
| 5 | T2 | Read PROD_QOH | 35 |
| 6 | T2 | $\text{PROD_QOH} = 35 - 30$ | |
| 7 | T2 | Write PROD_QOH | 5 |

TABLE
10.7

An Uncommitted Data Problem

| TIME | TRANSACTION | STEP | STORED VALUE |
|------|-------------|--|--------------|
| 1 | T1 | Read PROD_QOH | 35 |
| 2 | T1 | $\text{PROD_QOH} = 35 + 100$ | |
| 3 | T1 | Write PROD_QOH | 135 |
| 4 | T2 | Read PROD_QOH (Read uncommitted data) | 135 |
| 5 | T2 | $\text{PROD_QOH} = 135 - 30$ | |
| 6 | T1 | ***** ROLLBACK ***** | 35 |
| 7 | T2 | Write PROD_QOH | 105 |

Inconsistent Retrievals

- Inconsistent retrievals:
 - First transaction accesses data
 - Second transaction alters the data
 - First transaction accesses the data again
- Transaction might read some data before they are changed and other data after changed
- Yields inconsistent results

Inconsistent Retrievals (Cont...)

TABLE 10.8 Retrieval During Update

| TRANSACTION T1 | | TRANSACTION T2 | |
|----------------|---------------|----------------|--------------------------|
| SELECT | SUM(PROD_QOH) | UPDATE | PRODUCT |
| FROM | PRODUCT | SET | PROD_QOH = PROD_QOH + 10 |
| | | WHERE | PROD_CODE = '1546-QQ2' |
| | | UPDATE | PRODUCT |
| | | SET | PROD_QOH = PROD_QOH - 10 |
| | | WHERE | PROD_CODE = '1558-QW1' |
| | | COMMIT; | |

TABLE 10.9 Transaction Results: Data Entry Correction

| | BEFORE | AFTER |
|-----------|----------|----------------------------|
| PROD_CODE | PROD_QOH | PROD_QOH |
| 11QER/31 | 8 | 8 |
| 13-Q2/P2 | 32 | 32 |
| 1546-QQ2 | 15 | $(15 + 10) \rightarrow 25$ |
| 1558-QW1 | 23 | $(23 - 10) \rightarrow 13$ |
| 2232-QTY | 8 | 8 |
| 2232-QWE | 6 | 6 |
| Total | 92 | 92 |

Inconsistent Retrievals (Cont...)

TABLE
10.10

Inconsistent Retrievals

| TIME | TRANSACTION | ACTION | VALUE | TOTAL |
|------|-------------|---|-------|-------------|
| 1 | T1 | Read PROD_QOH for PROD_CODE = '11QER/31' | 8 | 8 |
| 2 | T1 | Read PROD_QOH for PROD_CODE = '13-Q2/P2' | 32 | 40 |
| 3 | T2 | Read PROD_QOH for PROD_CODE = '1546-QQ2' | 15 | |
| 4 | T2 | PROD_QOH = 15 + 10 | | |
| 5 | T2 | Write PROD_QOH for PROD_CODE = '1546-QQ2' | 25 | |
| 6 | T1 | Read PROD_QOH for PROD_CODE = '1546-QQ2' | 25 | (After) 65 |
| 7 | T1 | Read PROD_QOH for PROD_CODE = '1558-QW1' | 23 | (Before) 88 |
| 8 | T2 | Read PROD_QOH for PROD_CODE = '1558-QW1' | 23 | |
| 9 | T2 | PROD_QOH = 23 - 10 | | |
| 10 | T2 | Write PROD_QOH for PROD_CODE = '1558-QW1' | 13 | |
| 11 | T2 | ***** COMMIT ***** | | |
| 12 | T1 | Read PROD_QOH for PROD_CODE = '2232-QTY' | 8 | 96 |
| 13 | T1 | Read PROD_QOH for PROD_CODE = '2232-QWE' | 6 | 102 |

Scheduling Transactions

- Serial schedule: Schedule that does not interleave the actions of different transactions.
- Equivalent schedules: For any database state, the effect (on the set of objects in the database) of executing the first schedule is identical to the effect of executing the second schedule.
- Serializable schedule: A schedule that is equivalent to some serial execution of the transactions. (Note: If each transaction preserves consistency, every serializable schedule preserves consistency.)

The Scheduler


- As long as two transactions access unrelated data, there is no conflict in the execution (order is irrelevant to the final outcome)
- Special DBMS program
 - Purpose is to establish order of operations within which concurrent transactions are executed
- Interleaves execution of database operations:
 - Ensures **serializability**
 - Ensures **isolation**
- Serializable schedule
 - Interleaved execution of transactions yields same results **as some** serial execution

The Scheduler (Cont...)

- Bases its actions on concurrency control algorithms
- Ensures computer's central processing unit (CPU) is used efficiently
 - First-come first-served scheduling wastes processing time when CPU waits for READ or WRITE operation
- Facilitates data isolation to ensure that two transactions do not update **same data element** at same time

TABLE
10.11

READ/WRITE Conflict Scenarios: Conflicting Database Operations Matrix

| Operations  | TRANSACTIONS | | RESULT |
|---|--------------|-------|-------------|
| | T1 | T2 | |
| | Read | Read | No conflict |
| | Read | Write | Conflict |
| | Write | Read | Conflict |
| | Write | Write | Conflict |

For the same data unit

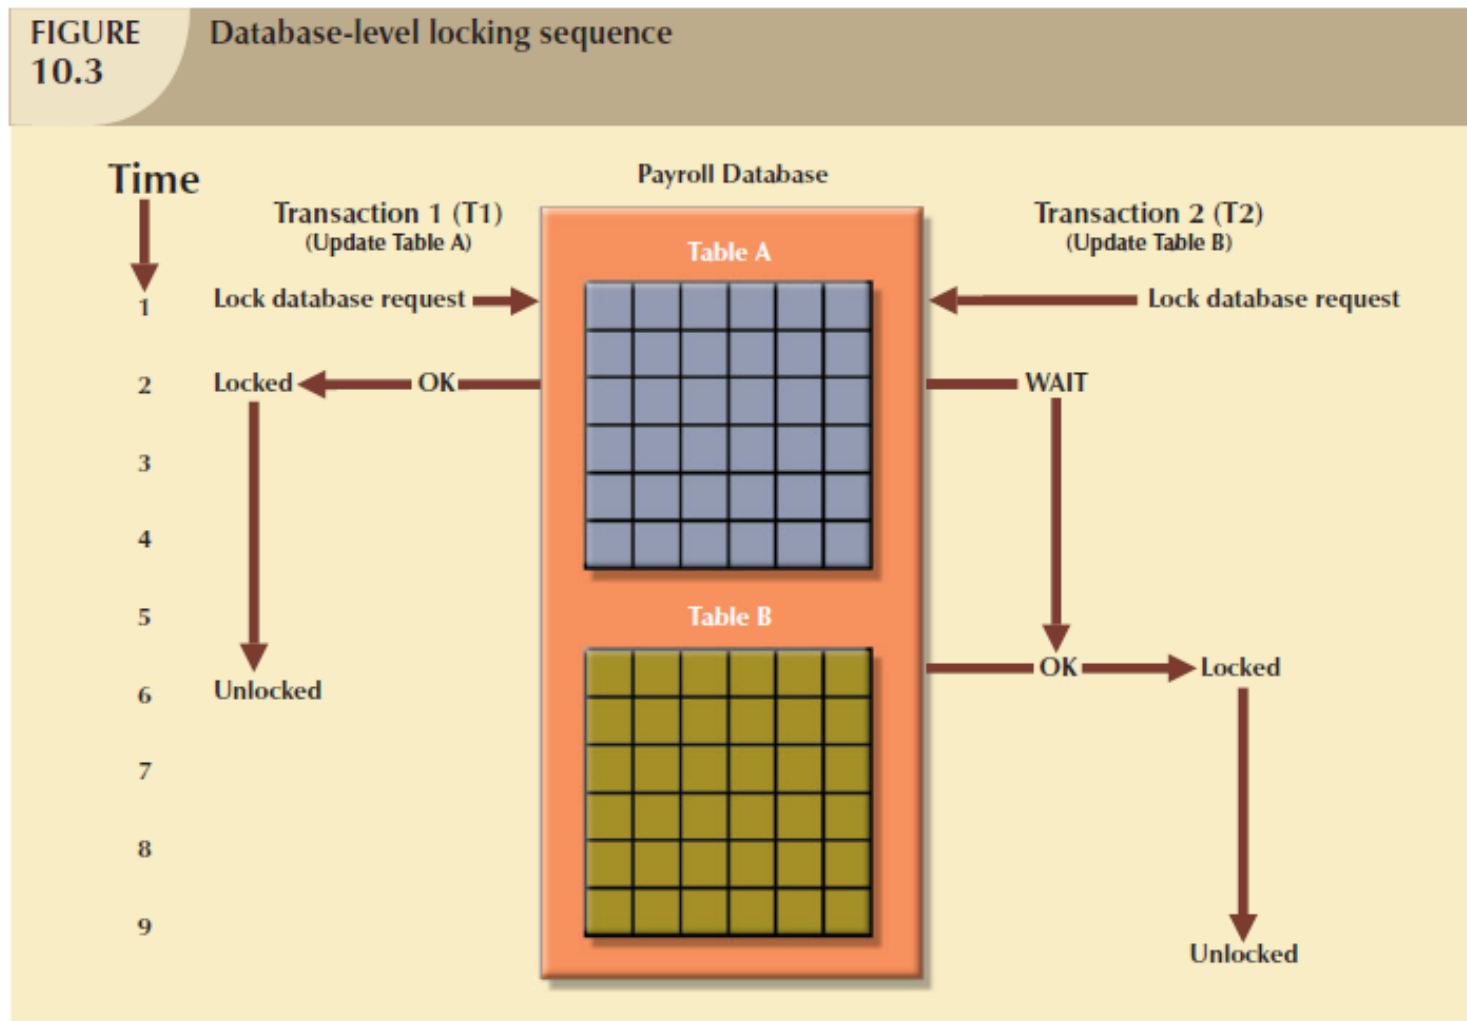
Concurrency Control with Locking Methods

- Lock
 - Guarantee exclusive use of a data item to a current transaction
 - Required to prevent another transaction from reading inconsistent data
- Lock Manager
 - Responsible for assigning and policing the locks used by transactions

Lock Granularity

- Indicates level of lock use
- Locking can take place at following levels:
 - Database: Entire database is locked
 - Table: Entire table is locked
 - Page: Entire diskpage is locked
 - Row:
 - Allows concurrent transactions to access different rows of same table, even if rows are located on same page
 - Field (attribute)
 - Allows concurrent transactions to access same row, as long as they require use of different fields (attributes) within that row

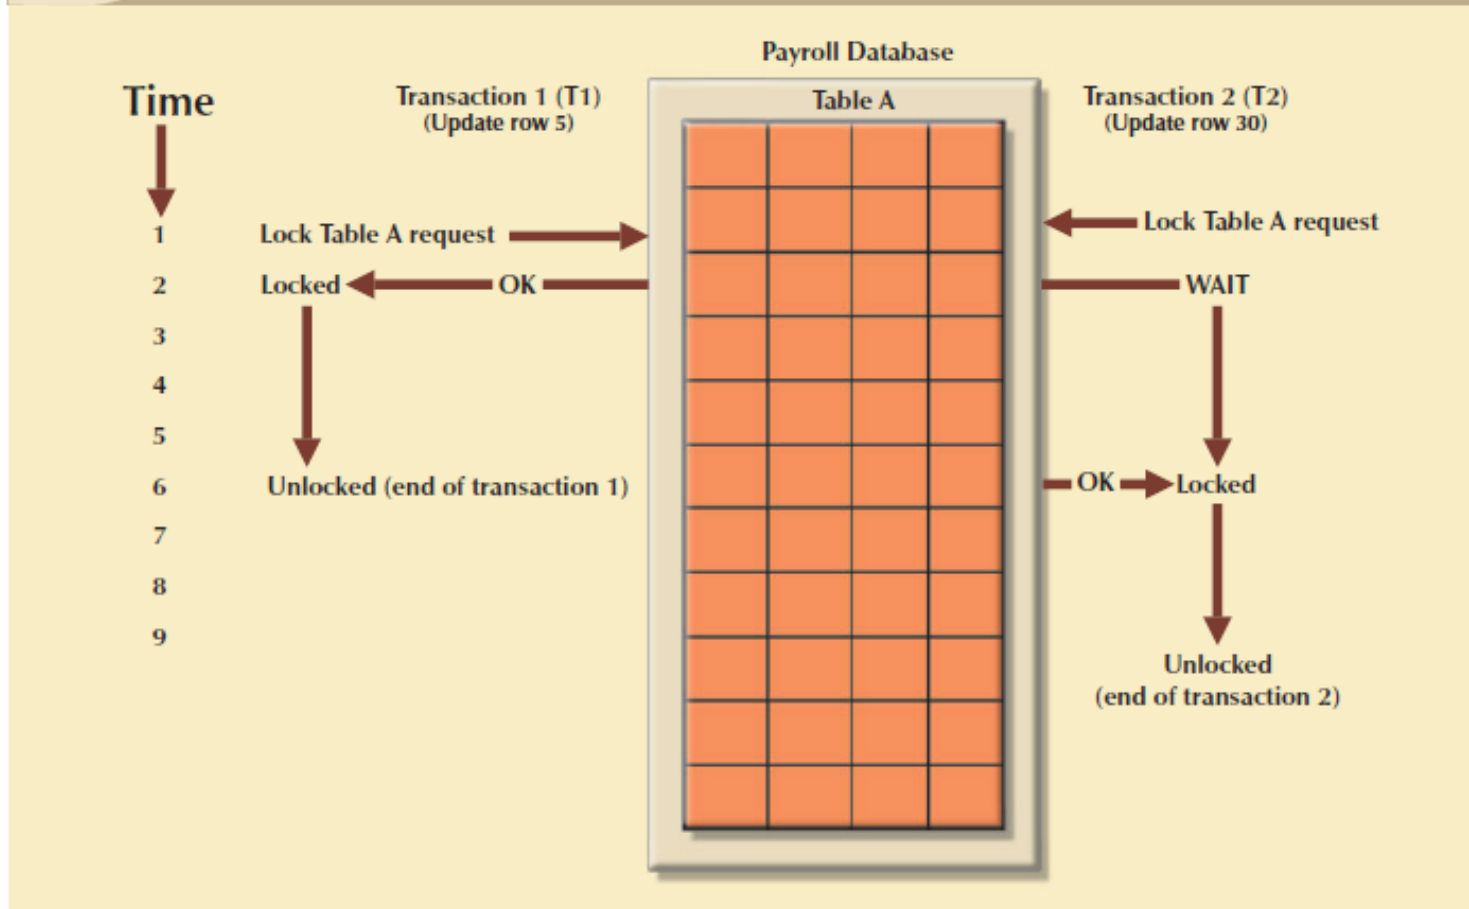
Database-level Locking Sequence



Example: Table-level Lock

FIGURE 10.4

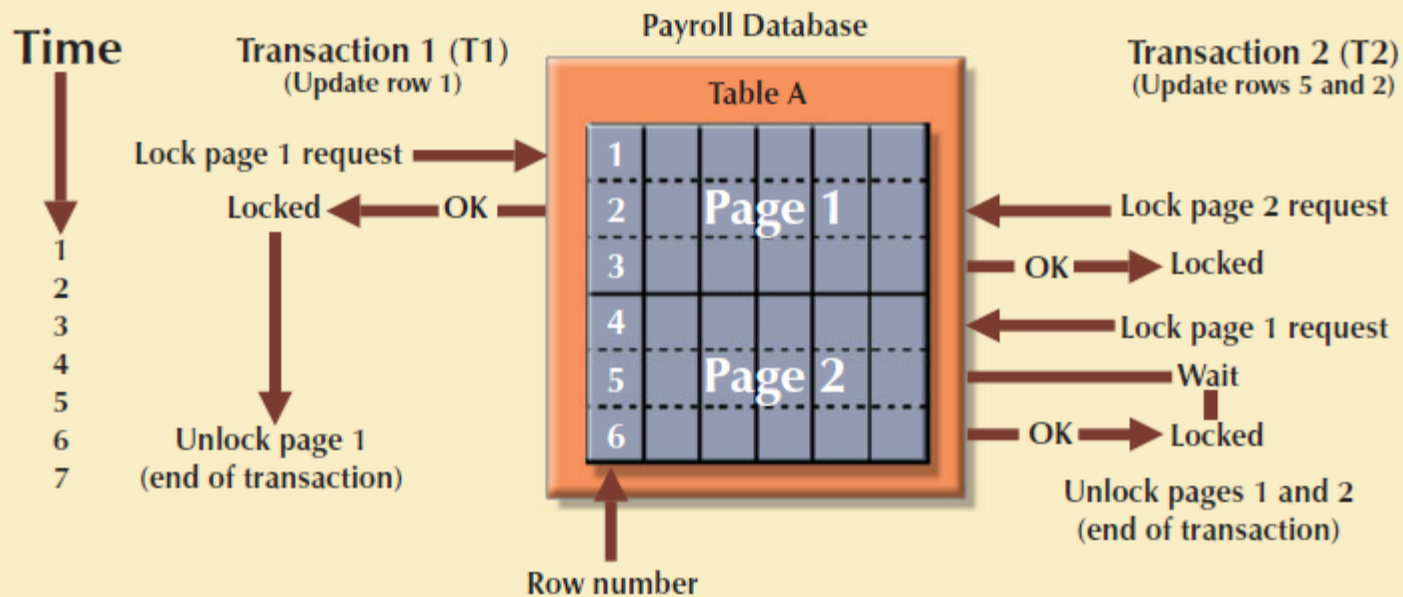
An example of a table-level lock



Example: Page-level Lock

FIGURE 10.5

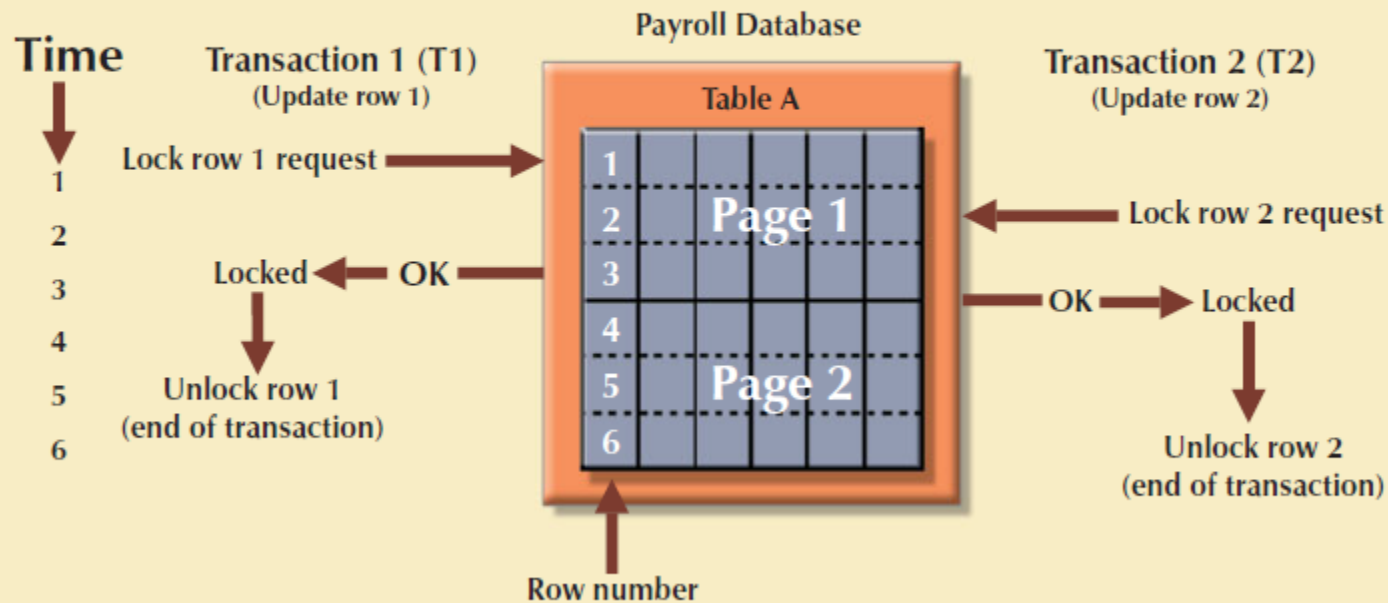
An example of a page-level lock



Example: Row-level Lock

FIGURE 10.6

An example of a row-level lock



Lock Types

- [Binary lock](#)
 - Two states: locked (1) or unlocked (0)
 - Every transaction required a lock and unlock operation for each accessed data item, which are automatically managed by the DBMS
- Exclusive lock
 - Access is specifically reserved for transaction that locked object
 - Mutual exclusive rule
 - Must be used when potential for conflict exists
- Shared lock
 - Concurrent transactions are granted read access on basis of a common lock

Locking Conflict Table

| Data Status \ Request | Not Locked | Share Locked | Exclusive Lock |
|-----------------------|-------------|--------------|----------------|
| Shared Lock | No Conflict | No Conflict | Conflict |
| Exclusive Lock | No Conflict | Conflict | Conflict |

Example: Binary Lock

TABLE
10.12

An Example of a Binary Lock

| TIME | TRANSACTION | STEP | STORED VALUE |
|------|-------------|--------------------|--------------|
| 1 | T1 | Lock PRODUCT | |
| 2 | T1 | Read PROD_QOH | 13 |
| 3 | T1 | PROD_QOH = 13 + 10 | |
| 4 | T1 | Write PROD_QOH | 23 |
| 5 | T1 | Unlock PRODUCT | |
| 6 | T2 | Lock PRODUCT | |
| 7 | T2 | Read PROD_QOH | 23 |
| 8 | T2 | PROD_QOH = 23 - 10 | |
| 9 | T2 | Write PROD_QOH | 13 |
| 10 | T2 | Unlock PRODUCT | |

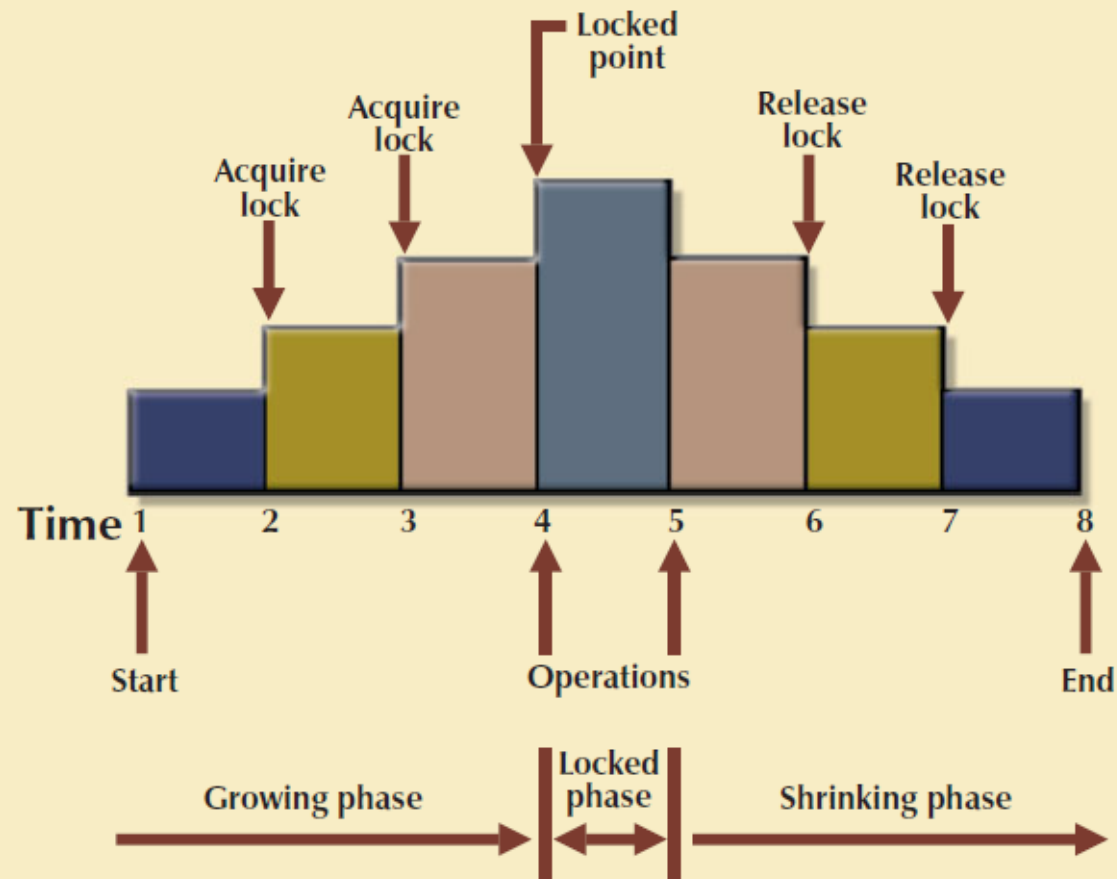
Two-Phase Locking to Ensure Serializability

- Defines how transactions acquire and relinquish locks
- Guarantees serializability, but does not prevent deadlocks
 - Growing phase
 - Transaction acquires all required locks **without unlocking** any data
 - Shrinking phase
 - Transaction releases all locks and **cannot obtain any new lock**
- Governed by the following rules:
 - Two transactions cannot have conflicting locks
 - No unlock operation can precede a lock operation in the same transaction
 - No data are affected until all locks are obtained
 - That is, until transaction in in its locked point

Two-Phase Locking Protocol

FIGURE 10.7

Two-phase locking protocol



Deadlocks

- Condition that occurs when two transactions wait for each other to unlock data
- Possible only if one of the transactions wants to obtain an exclusive lock on a data item
 - No deadlock condition can exist among shared locks
- Three techniques to control deadlock:
 - Prevention: Abort & reschedule transaction if deadlock possible
 - Detection: Test database for deadlocks; abort victim, leave rest
 - Avoidance
- Choice of deadlock control method depends on database environment
 - Low probability of deadlock, detection recommended
 - High probability of deadlock, prevention recommended

How Deadlock Condition is Created?

TABLE
10.13

How a Deadlock Condition Is Created

| TIME | TRANSACTION | REPLY | LOCK STATUS | |
|------|-------------|-------|-------------|----------|
| 0 | | | Data X | Data Y |
| 1 | T1:LOCK(X) | OK | Unlocked | Unlocked |
| 2 | T2: LOCK(Y) | OK | Locked | Unlocked |
| 3 | T1:LOCK(Y) | WAIT | Locked | Locked |
| 4 | T2:LOCK(X) | WAIT | Locked | Locked |
| 5 | T1:LOCK(Y) | WAIT | Locked | Locked |
| 6 | T2:LOCK(X) | WAIT | Locked | Locked |
| 7 | T1:LOCK(Y) | WAIT | Locked | Locked |
| 8 | T2:LOCK(X) | WAIT | Locked | Locked |
| 9 | T1:LOCK(Y) | WAIT | Locked | Locked |
| ... | | | | |
| ... | | | | |
| ... | | | | |
| ... | | | | |

Deadlock



Concurrency Control with Time Stamping Methods

- Assign global, unique time stamp to each transaction
- Produces explicit order in which transactions are submitted to DBMS
- Properties of time stamp
 - **Uniqueness**
 - Ensures that no equal time stamp values can exist
 - **Monotonicity**
 - Ensures that time stamp values always increase
- Disadvantages
 - Each value in database needs two fields: a) last time field was read, and b) last update
 - Increases memory needs and processing overhead

Wait/Die and Wound/Wait Schemes

- Wait/die
 - Older requesting transaction **waits** and
 - Younger requesting transaction is **rolled back** and rescheduled
- Wound/wait
 - Older requesting transaction **preempts** (rolls back) younger transaction and reschedules it
 - Younger requesting transaction **waits**

Wait/Die and Wound/Wait Schemes (Cont...)

TABLE
10.14

Wait/Die and Wound/Wait Concurrency Control Schemes

| TRANSACTION REQUESTING LOCK | TRANSACTION OWNING LOCK | WAIT/DIE SCHEME | WOUND/WAIT SCHEME |
|-----------------------------|-------------------------|--|---|
| T1 (11548789) | T2 (19562545) | <ul style="list-style-type: none">• T1 waits until T2 is completed and T2 releases its locks. | <ul style="list-style-type: none">• T1 preempts (rolls back) T2.• T2 is rescheduled using the same time stamp. |
| T2 (19562545) | T1 (11548789) | <ul style="list-style-type: none">• T2 dies (rolls back)• T2 is rescheduled using the same time stamp | <ul style="list-style-type: none">• T2 waits until T1 is completed and T1 releases its locks. |

Concurrency Control with Optimistic Methods

- Optimistic approach
 - Based on assumption that **majority** of database operations **do not conflict**
 - Does not require locking or time stamping techniques
 - Transaction is **executed without restrictions until it is committed**
 - Phases: read, validation, and write
 - Good for read/query database systems requiring few update transactions; Poor for heavily used DBMS environment

Reasons for a Crash

- System Crash
- Transaction or System Error
- Local Error or Exception
- Concurrency Control
- Disk Failure
- Catastrophe

Crash Recovery

- Recovery Manager
 - When DBMS is restarted after crashes, the recovery manager must bring the database to a consistent state
 - Ensures transaction atomicity and durability
 - Undo actions of transactions that do not commit
 - Redo actions of committed transactions during system failures and media failures (corrupted disk).
- Recovery Manager maintains log information during normal execution of transactions for use during crash recovery

Database Recovery Management

- Restores database to previous consistent state
- Based on atomic transaction property
 - All portions of transaction treated as single logical unit of work
 - All operations applied and completed to produce consistent database
- If transaction operation cannot be completed
 - Transaction aborted
 - Changes to database rolled back (**undone**)

Concept that Affect Transaction Recovery

- **Write-ahead-log protocol:** ensures transaction logs are written before data is updated
- **Redundant transaction logs:** ensure physical disk failure will not impair ability to recover
- **Buffers:** temporary storage areas in primary memory
- **Checkpoints:** operations in which DBMS writes all its updated buffers to disk

Transaction Recovery

- Make use of deferred-write and write-through techniques
- **Deferred-write technique**
 - Transaction operations do not immediately update physical database
 - Only transaction log is updated
 - Database is physically updated only after transaction reaches its commit point using transaction log information

Transaction Recovery (Cont...)

- Recovery process for deferred-write:
 - Identify last checkpoint
 - If transaction **committed before** checkpoint
 - Do nothing
 - If transaction committed after checkpoint
 - Use transaction log **to redo** the transaction
 - If transaction had ROLLBACK operation or was left active
 - **Do nothing because no updates were made**

Transaction Recovery (Cont...)

- **Write-through technique**
 - Database is **immediately updated** by transaction operations during transaction's execution
- Recovery process for write-through
 - Identify last checkpoint
 - If transaction committed before checkpoint
 - Do nothing
 - If transaction committed after checkpoint
 - DBMS redoes the transaction using "after" values
 - If transaction had ROLLBACK operation or was left active
 - **Uses the before value in the transaction log records to ROLLBACK (undo)**

Transaction Recovery (Cont...)

TABLE
10.15

A Transaction Log for Transaction Recovery Examples

| TRL ID | TRX NUM | PREV PTR | NEXT PTR | OPERATION | TABLE | ROW ID | ATTRIBUTE | BEFORE VALUE | AFTER VALUE |
|--------------------------|---------|----------|----------|------------|-------------------------|----------|--------------|--------------|-------------------------|
| 341 | 101 | Null | 352 | START | ****Start Transaction | | | | |
| 352 | 101 | 341 | 363 | UPDATE | PRODUCT | 54778-2T | PROD_QOH | 45 | 43 |
| 363 | 101 | 352 | 365 | UPDATE | CUSTOMER | 10011 | CUST_BALANCE | 615.73 | 675.62 |
| 365 | 101 | 363 | Null | COMMIT | **** End of Transaction | | | | |
| 397 | 106 | Null | 405 | START | ****Start Transaction | | | | |
| 405 | 106 | 397 | 415 | INSERT | INVOICE | 1009 | | | 1009,10016, ... |
| 415 | 106 | 405 | 419 | INSERT | LINE | 1009,1 | | | 1009,1, 89-WRE-Q,1, ... |
| 419 | 106 | 415 | 427 | UPDATE | PRODUCT | 89-WRE-Q | PROD_QOH | 12 | 11 |
| 423 | | | | CHECKPOINT | | | | | |
| 427 | 106 | 419 | 431 | UPDATE | CUSTOMER | 10016 | CUST_BALANCE | 0.00 | 277.55 |
| 431 | 106 | 427 | 457 | INSERT | ACCT_TRANSACTION | 10007 | | | 1007,18-JAN-2008, ... |
| 457 | 106 | 431 | Null | COMMIT | **** End of Transaction | | | | |
| 521 | 155 | Null | 525 | START | ****Start Transaction | | | | |
| 525 | 155 | 521 | 528 | UPDATE | PRODUCT | 2232/QWE | PROD_QOH | 6 | 26 |
| 528 | 155 | 525 | Null | COMMIT | **** End of Transaction | | | | |
| ***** C *R*A* S* H ***** | | | | | | | | | |

Summary

- Transaction: sequence of database operations that access database
 - Logical unit of work
 - No portion of transaction can exit by itself
 - Five main properties: atomicity, consistency, isolation, durability, and serializability
- COMMIT saves changes to disk
- ROLLBACK restores previous database state
- SQL transactions are formed by several SQL statements or database requests

Summary (Cont...)

- Transaction log keeps track of all transactions that modify database
- Concurrency control coordinates simultaneous execution of transactions
- Scheduler establishes order in which concurrent transaction operations are executed
- Lock guarantee unique access to data item by transaction
- Two types of locks: binary locks and shared/exclusive locks

Summary (Cont...)

- Serializability of schedules is guaranteed through the use of two-phase locking
- Deadlock: when two or more transactions wait indefinitely for each other to release lock
- Three deadlock control techniques: prevention, detection, and avoidance
- Time stamping methods assign unique time stamp to each transaction
 - Schedules execution of conflicting transactions in time stamp order

Summary (Cont...)

- Optimistic methods assume the majority of database transactions do not conflict
 - Transactions are executed concurrently, using private copies of the data
- Database recovery restores database from given state to previous consistent state