

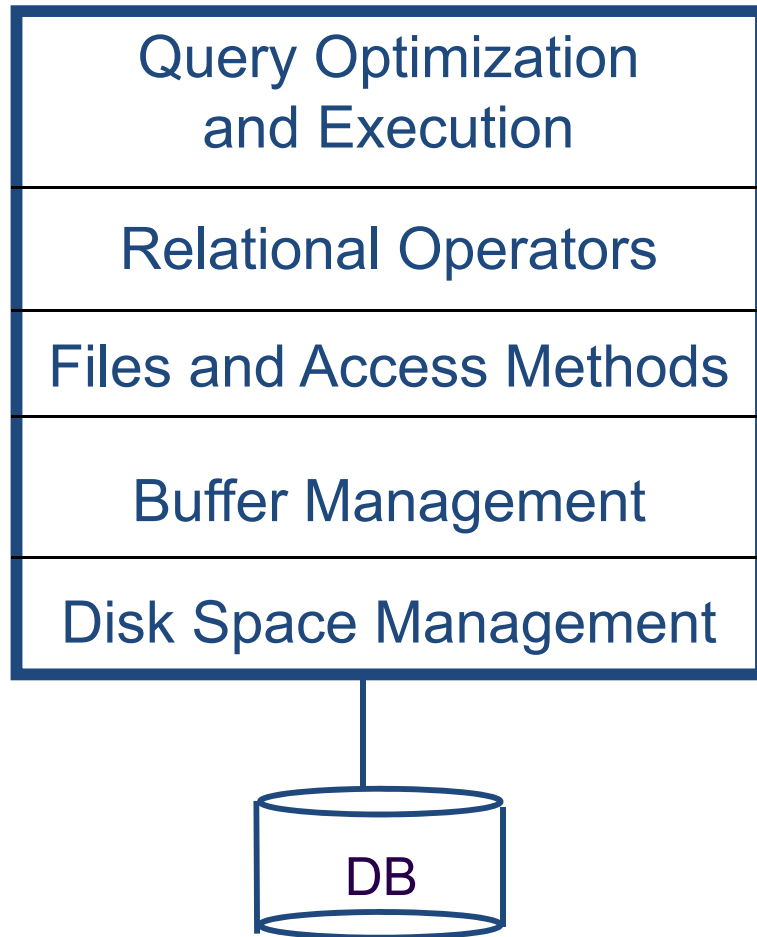
# Database Management System

Sumayyea Salahuddin (Lecturer)  
Dept. of Computer Systems Eng.  
UET Peshawar

# Overview

- Database **transactions** and their properties
- What **concurrency control** is and what role it plays in maintaining the database's integrity
- What **locking** methods are and how they work
- How **stamping** methods are used for concurrency control
- How **optimistic methods** are used for concurrency control
- How database **recovery management** is used to maintain database integrity

# DBMS Structure



These layers must consider concurrency control and recovery (**Transaction, Lock, Recovery Managers**)

# What is Transaction?

- Logical unit of work that must be either entirely completed or aborted
- Successful transaction changes database from one consistent state to another
  - One in which all data integrity constraints are satisfied
- Most real-world database transactions are formed by two or more database requests
  - Equivalent of a single SQL statement in an application program or transaction
    - if a transaction is composed of two UPDATE statements and one INSERT statement, the transaction uses three database requests. In turn, each database request generates several input/output (I/O) operations that read from or write to physical storage media.

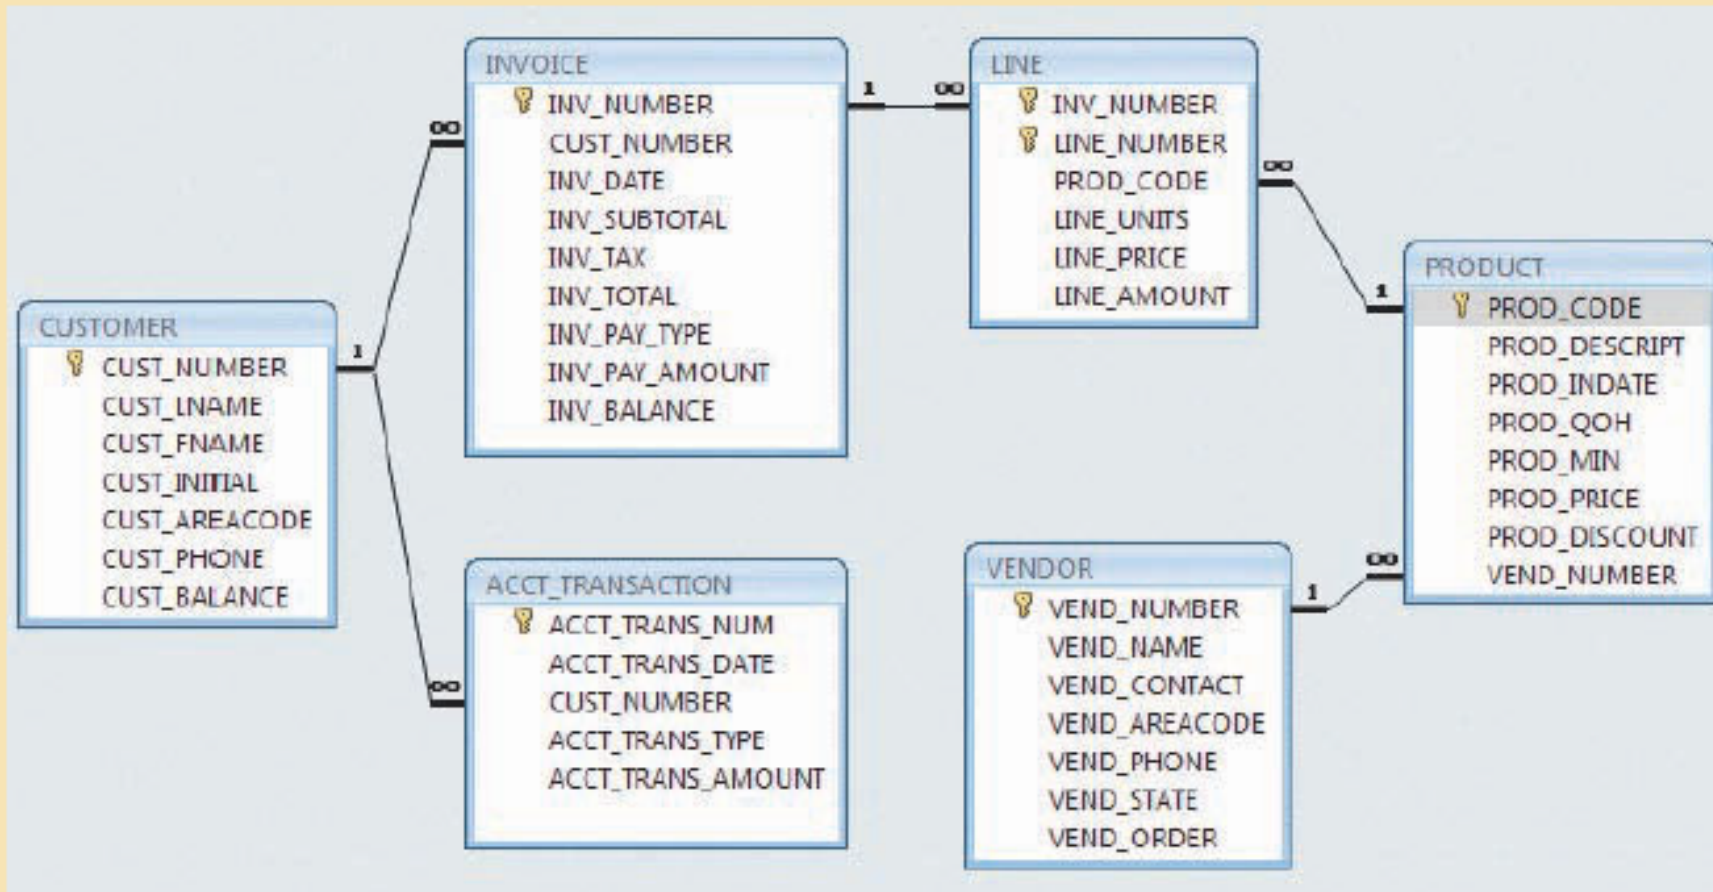
# What is Transaction? (Cont...)

- A transaction is any action that reads from and/or writes to a database & may consist of:
  - Simple SELECT statement to generate a list of table contents
  - Series of related UPDATE statements to change values of attributes in various tables
  - Series of INSERT statements to add rows to one or more tables
  - Combination of SELECT, UPDATE, and INSERT statements
- Example in next slide

# Sample Relational Database

FIGURE  
10.1

The Ch10\_SaleCo database relational diagram



# Example: Sales Transaction

- START TRANSACTION;
- INSERT INTO INVOICE VALUES (1009, 10016, '18-Jan-2008', 256.99, 20.56, 277.55, 'cred', 0.00, 277.55);
- INSERT INTO LINE VALUES (1009, 1, '89-WRE-Q', 1, 256.99, 256.99);
- UPDATE PRODUCT SET PROD\_QOH = PROD\_QOH - 1 WHERE PROD\_CODE = '89-WRE-Q';
- UPDATE CUSTOMER SET CUST\_BALANCE = CUST\_BALANCE + 277.55 WHERE CUST\_NUMBER = 10016;
- INSERT INTO ACCT\_TRANSACTION VALUES (10007, '18-Jan-08', 10016, 'charge', 277.55);
- COMMIT/ROLLBACK;

# Transaction Results

Table name: INVOICE

INV_NUMBER	CUST_NUMBER	INV_DATE	INV_SUBTOTAL	INV_TAX	INV_TOTAL	INV_PAY_TYPE	INV_PAY_AMOUNT	INV_BALANCE
1001	10014	16-Jan-08	54.02	4.30	59.31	cc	59.31	0.00
1002	10011	16-Jan-08	9.98	0.80	10.78	cash	10.78	0.00
1003	10012	10-Jan-00	270.70	21.00	292.00	cc	292.00	0.00
1004	10011	17-Jan-08	34.87	2.79	37.66	cc	37.66	0.00
1005	10010	17-Jan-08	70.44	5.04	75.00	cc	75.00	0.00
1006	10014	17-Jan-08	397.83	31.83	429.66	cred	100.00	329.66
1007	10015	17-Jan-08	34.97	2.80	37.77	chd	37.77	0.00
1008	10011	17-Jan-08	1033.08	82.65	1115.73	cred	500.00	615.73
1009	10016	18-Jan-08	256.99	20.58	277.55	cred	0.00	277.55

Table name: LINE

INV_NUMBER	LINE_NUMBER	PROD_CODE	LINE_UNITS	LINE_PRICE	LINE_AMOUNT
1001	1	13-Q2/P2	3	14.00	44.07
1001	2	23109-HB	1	9.95	9.95
1002	1	54778-2T	2	4.00	9.98
1003	1	2338106D	4	38.95	155.80
1003	2	1540-QQ2	1	39.95	39.95
1003	3	13-Q2/P2	6	14.00	74.06
1004	1	54778-2T	3	4.99	14.97
1004	2	23109-HB	2	9.95	19.90
1005	1	PVC23DRT	12	5.87	70.44
1006	1	SM-16277	3	6.00	20.07
1006	2	2232QTY	1	109.92	109.92
1006	3	23109-HB	1	0.05	0.05
1006	4	89-WIRE-Q	1	256.99	256.99
1007	1	13-Q2/P2	2	14.99	29.90
1007	2	54778-2T	1	4.99	4.99
1008	1	PVC23DRT	5	5.07	29.35
1008	2	VR31T3	4	119.95	479.80
1008	3	23109-HB	1	9.95	9.95
1008	4	89-WIRE-Q	2	256.00	513.08
1009	1	89-WIRE-Q	1	256.99	256.99

Table name: PRODUCT

PROD_CODE	PROD_DESCRIPTION	PROD_INDATE	PROD_QOH	PROD_MIN	PROD_PRICE	PROD_DISCOUNT	VEND_NUMBER
10ER/31	Power painter, 15 psi, 3-nozzle	03-Nov-07	8	5	100.00	0.00	25505
13-Q2/P2	7.25-in. pwr. saw blade	13-Dec-07	32	15	14.99	0.05	21344
14-Q1/L3	9.00-in. pwr. saw blade	13-Nov-07	10	12	17.49	0.00	21344
1516-QQ2	Hnd. cloth, 14 in., 2x50	15-Jan-08	15	8	39.95	0.00	23119
1550-QVM	Hnd. cloth, 12 in., 3x50	15-Jan-08	20	5	40.99	0.00	23119
2232QTY	B&D jigsaw, 12 in. blade	30-Dec-07	8	5	109.92	0.05	24288
2232QMF	B&D jigsaw, 8 in. blade	24-Dec-07	6	5	99.87	0.05	24288
2238GPD	B&D cordless drill, 1/2 in.	20-Jan-08	12	5	38.95	0.05	25505
23109-HB	Crew hammer	20-Jan-08	23	10	9.95	0.10	21225
23114-AA	Sledge hammer, 12 lb.	02-Jan-08	8	5	14.40	0.05	21225
54778-2T	Rat-tail file, 1/8 in. fine	15-Dec-07	43	20	4.99	0.00	21344
89-WIRE-Q	Licut chain saw, 10 in.	07-Jan-00	11	5	250.99	0.05	24200
PVC23DRT	PVC pipe, 3.5 in., 8 ft	06-Jan-08	188	75	5.87	0.00	21225
SM-16277	1.25-in. metal screw, 25	01-Mar-00	172	75	0.99	0.00	21225
SM-23116	2.5-in. wd. screw, 50	24-Feb-08	237	100	8.45	0.00	21231
VR31T3	Steel netting, 4x8x1.6", .5" mesh	17-Jan-08	18	5	119.95	0.10	25595

Table name: CUSTOMER

CUST_NUMBER	CUST_LNAME	CUST_FNAME	CUST_INITIAL	CUST_AREACODE	CUST_PHONE	CUST_BALANCE
10010	Ramos	Alfred	A	615	844-2573	0.00
10011	Dumas	Lucas	K	713	894-1298	615.73
10012	Gmth	Kathy	W	616	804-2296	0.00
10013	Okrowski	Paul	F	815	894-2180	0.00
10014	Orlano	Myron		615	222-1872	0.00
10015	O'Brien	Amy	B	713	442-3381	0.00
10016	Brown	James	Q	615	207-1228	277.55
10017	Williams	George		815	290-2558	0.00
10018	Farnes	Anne	Q	713	382-7185	0.00
10019	Smith	Olette	K	615	297-3809	0.00

Table name: ACCT TRANSACTION

ACCT_TRANS_NUM	ACCT_TRANS_DATE	CUST_NUMBER	ACCT_TRANS_TYPE	ACCT_TRANS_AMOUNT
10003	17-Jan-08	10014	charge	329.66
10004	17-Jan-08	10011	charge	615.73
10005	23-Jan-08	10014	payment	329.66
10007	18-Jan-08	10016	charge	277.55



# Evaluating Transaction Results

- Not all transactions update database
- SQL code represents a transaction because database was accessed
- Improper or incomplete transactions can have devastating effect on database integrity
  - Some DBMSs provide means by which user can define enforceable constraints
  - Other integrity rules are enforced automatically by the DBMS
    - Primary key integrity
    - Referential integrity
    - Entity integrity
- **No semantic/logical checking**
- Transactions that violate integrity constraints are aborted

# Transaction Properties

- **Atomicity**
  - All operations of a transaction must be completed
- **Consistency**
  - Permanence of database's consistent state
- **Isolation**
  - Data used during transaction cannot be used by second transaction until the first is completed
- **Durability**
  - Once transactions are committed, they cannot be undone
- **Serializability**
  - Concurrent execution of several transactions yields consistent results

Multituser databases subject to multiple concurrent transactions

# Transaction Management with SQL

- ANSI has defined standards that govern SQL database transactions
- Transaction support is provided by two SQL statements: **COMMIT** and **ROLLBACK**
- Transaction sequence must continue until:
  - COMMIT statement is reached
  - ROLLBACK statement is reached
  - End of program is reached (Equivalent to COMMIT)
  - Program is abnormally terminated (Equivalent to ROLLBACK)

# Transaction & Concurrent Execution

- **Transaction Manager** controls the execution of transactions.
- User program may carry out many operations on the data retrieved from the database, but the DBMS is only concerned about what data is read/written from/to the database.
- Concurrent execution of multiple transactions is essential for good performance.
  - Disk is the bottleneck (slow, frequently used)
  - Must keep CPU busy w/many queries
  - Better response time

# The Transaction Log

- Keeps track of all transactions that **update the database**
- Useful for recovering database
- **Transaction log** stores:
  - A record for the beginning of transaction
  - For each transaction component:
    - Type of operation being performed (update, delete, insert)
    - Names of objects affected by transaction
    - "Before" and "after" values for updated fields
    - Pointers to previous and next transaction log entries for the same transaction
  - Ending (COMMIT) of the transaction

# Sample Transaction Log

TABLE  
10.1

A Transaction Log

TRL_ID	TRX_NUM	PREV_PTR	NEXT_PTR	OPERATION	TABLE	ROW ID	ATTRIBUTE	BEFORE VALUE	AFTER VALUE
341	101	Null	352	START	****Start Transaction				
352	101	341	363	UPDATE	PRODUCT	1558-QW1	PROD_QOH	25	23
363	101	352	365	UPDATE	CUSTOMER	10011	CUST_BALANCE	525.75	615.73
365	101	363	Null	COMMIT	**** End of Transaction				



TRL\_ID = Transaction log record ID

PTR = Pointer to a transaction log record ID

TRX\_NUM = Transaction number

(Note: The transaction number is automatically assigned by the DBMS.)

START TRANSACTION;

UPDATE PRODUCT SET PROD\_QOH = 23 WHERE PROD\_ID = '1558-QW1';

UPDATE CUSTOMER SET CUST\_BALANCE = 615.73 WHERE CUST\_ID = 10011;

COMMIT;

# Concurrency Control

- Coordination of simultaneous transaction execution in a multiprocessing database
- Objective is to ensure serializability of transactions in a multiuser environment
- Simultaneous execution of transactions over a shared database can create several data integrity and consistency problems
  - Lost updates
  - Uncommitted data
  - Inconsistent retrievals

# Lost Updates

- **Lost update** problem:
  - Two concurrent transactions update same data element
  - One of the update is lost
    - Overwritten by the other transaction

TABLE  
10.2

Two Concurrent Transactions to Update QOH

TRANSACTION	COMPUTATION
T1: Purchase 100 units	$\text{PROD\_QOH} = \text{PROD\_QOH} + 100$
T2: Sell 30 units	$\text{PROD\_QOH} = \text{PROD\_QOH} - 30$



# Lost Updates (Cont...)

TABLE  
10.3

Serial Execution of Two Transactions

TIME	TRANSACTION	STEP	STORED VALUE
1	T1	Read PROD_QOH	35
2	T1	$\text{PROD\_QOH} = 35 + 100$	
3	T1	Write PROD_QOH	135
4	T2	Read PROD_QOH	135
5	T2	$\text{PROD\_QOH} = 135 - 30$	
6	T2	Write PROD_QOH	105

TABLE  
10.4

Lost Updates

TIME	TRANSACTION	STEP	STORED VALUE
1	T1	Read PROD_QOH	35
2	T2	Read PROD_QOH	35
3	T1	$\text{PROD\_QOH} = 35 + 100$	
4	T2	$\text{PROD\_QOH} = 35 - 30$	
5	T1	Write PROD_QOH (Lost update)	135
6	T2	Write PROD_QOH	5

# Uncommitted Data

- **Uncommitted data phenomenon:**
  - Two transaction executed concurrently
  - First transaction rolled back after second already accessed uncommitted data

TABLE  
10.5

Transactions Creating Uncommitted Data Problem

TRANSACTION	COMPUTATION
T1: Purchase 100 units	PROD_QOH = PROD_QOH + 100 (Rolled back)
T2: Sell 30 units	PROD_QOH = PROD_QOH 30

# Uncommitted Data (Cont...)

TABLE  
10.6

Correct Execution of Two Transactions

TIME	TRANSACTION	STEP	STORED VALUE
1	T1	Read PROD_QOH	35
2	T1	$\text{PROD\_QOH} = 35 + 100$	
3	T1	Write PROD_QOH	135
4	T1	*****ROLLBACK*****	35
5	T2	Read PROD_QOH	35
6	T2	$\text{PROD\_QOH} = 35 - 30$	
7	T2	Write PROD_QOH	5

TABLE  
10.7

An Uncommitted Data Problem

TIME	TRANSACTION	STEP	STORED VALUE
1	T1	Read PROD_QOH	35
2	T1	$\text{PROD\_QOH} = 35 + 100$	
3	T1	Write PROD_QOH	135
4	T2	Read PROD_QOH (Read uncommitted data)	135
5	T2	$\text{PROD\_QOH} = 135 - 30$	
6	T1	*****ROLLBACK*****	35
7	T2	Write PROD_QOH	105

# Inconsistent Retrievals

- Inconsistent retrievals:
  - First transaction accesses data
  - Second transaction alters the data
  - First transaction accesses the data again
- Transaction might read some data before they are changed and other data after changed
- Yields inconsistent results

# Inconsistent Retrievals (Cont...)

TABLE  
10.8

Retrieval During Update

TRANSACTION T1		TRANSACTION T2	
SELECT	SUM(PROD_QOH)	UPDATE	PRODUCT
FROM	PRODUCT	SET	PROD_QOH = PROD_QOH + 10
		WHERE	PROD_CODE = '1546-QQ2'
		UPDATE	PRODUCT
		SET	PROD_QOH = PROD_QOH - 10
		WHERE	PROD_CODE = '1558-QW1'
		COMMIT;	

TABLE  
10.9

Transaction Results: Data Entry Correction

	BEFORE	AFTER
PROD_CODE	PROD_QOH	PROD_QOH
11QER/31	8	8
13-Q2/P2	32	32
1546-QQ2	15	$(15 + 10) \rightarrow 25$
1558-QW1	23	$(23 - 10) \rightarrow 13$
2232-QTY	8	8
2232-QWE	6	6
Total	92	92

# Inconsistent Retrievals (Cont...)

TABLE  
10.10

Inconsistent Retrievals

TIME	TRANSACTION	ACTION	VALUE	TOTAL
1	T1	Read PROD_QOH for PROD_CODE = '11QER/31'	8	8
2	T1	Read PROD_QOH for PROD_CODE = '13-Q2/P2'	32	40
3	T2	Read PROD_QOH for PROD_CODE = '1546-QQ2'	15	
4	T2	PROD_QOH = 15 + 10		
5	T2	Write PROD_QOH for PROD_CODE = '1546-QQ2'	25	
6	T1	Read PROD_QOH for PROD_CODE = '1546-QQ2'	25	(After) 65
7	T1	Read PROD_QOH for PROD_CODE = '1558-QW1'	23	(Before) 88
8	T2	Read PROD_QOH for PROD_CODE = '1558-QW1'	23	
9	T2	PROD_QOH = 23 - 10		
10	T2	Write PROD_QOH for PROD_CODE = '1558-QW1'	13	
11	T2	***** COMMIT *****		
12	T1	Read PROD_QOH for PROD_CODE = '2232-QTY'	8	96
13	T1	Read PROD_QOH for PROD_CODE = '2232-QWE'	6	102

# Scheduling Transactions

- Serial schedule: Schedule that does not interleave the actions of different transactions.
- Equivalent schedules: For any database state, the effect (on the set of objects in the database) of executing the first schedule is identical to the effect of executing the second schedule.
- Serializable schedule: A schedule that is equivalent to some serial execution of the transactions. (Note: If each transaction preserves consistency, every serializable schedule preserves consistency.)

# The Scheduler

- As long as two transactions access unrelated data, there is no conflict in the execution (order is irrelevant to the final outcome)
- Special DBMS program
  - Purpose is to establish order of operations within which concurrent transactions are executed
- Interleaves execution of database operations:
  - Ensures **serializability**
  - Ensures **isolation**
- Serializable schedule
  - Interleaved execution of transactions yields same results **as some** serial execution




# The Scheduler (Cont...)

- Bases its actions on concurrency control algorithms
- Ensures computer's central processing unit (CPU) is used efficiently
  - First-come first-served scheduling wastes processing time when CPU waits for READ or WRITE operation
- Facilitates data isolation to ensure that two transactions do not update **same data element** at same time

TABLE  
10.11

READ/WRITE Conflict Scenarios: Conflicting Database Operations Matrix

Operations 	TRANSACTIONS		RESULT
	T1	T2	
	Read	Read	No conflict
	Read	Write	Conflict
	Write	Read	Conflict
	Write	Write	Conflict

For the same data unit

# Concurrency Control with Locking Methods


- Lock
  - Guarantee exclusive use of a data item to a current transaction
  - Required to prevent another transaction from reading inconsistent data
- Lock Manager
  - Responsible for assigning and policing the locks used by transactions

# Lock Granularity

- Indicates level of lock use
- Locking can take place at following levels:
  - Database: Entire database is locked
  - Table: Entire table is locked
  - Page: Entire diskpage is locked
  - Row:
    - Allows concurrent transactions to access different rows of same table, even if rows are located on same page
  - Field (attribute)
    - Allows concurrent transactions to access same row, as long as they require use of different fields (attributes) within that row

# Database-level Locking Sequence

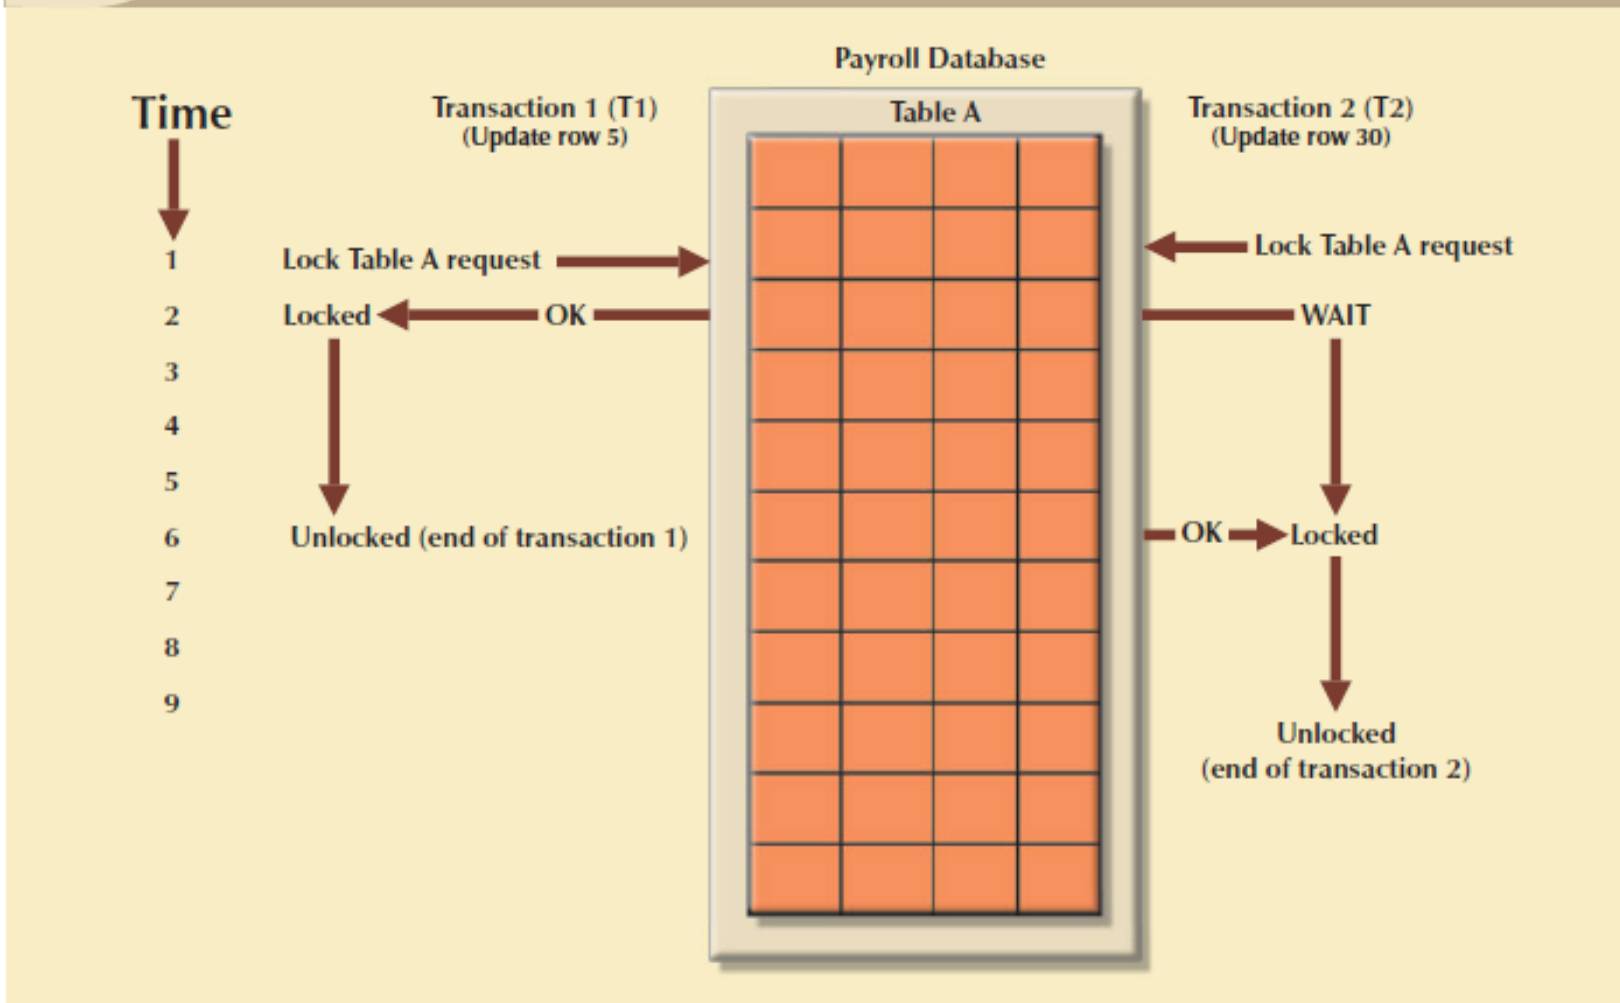
FIGURE 10.3 Database-level locking sequence



# Example: Table-level Lock

FIGURE 10.4

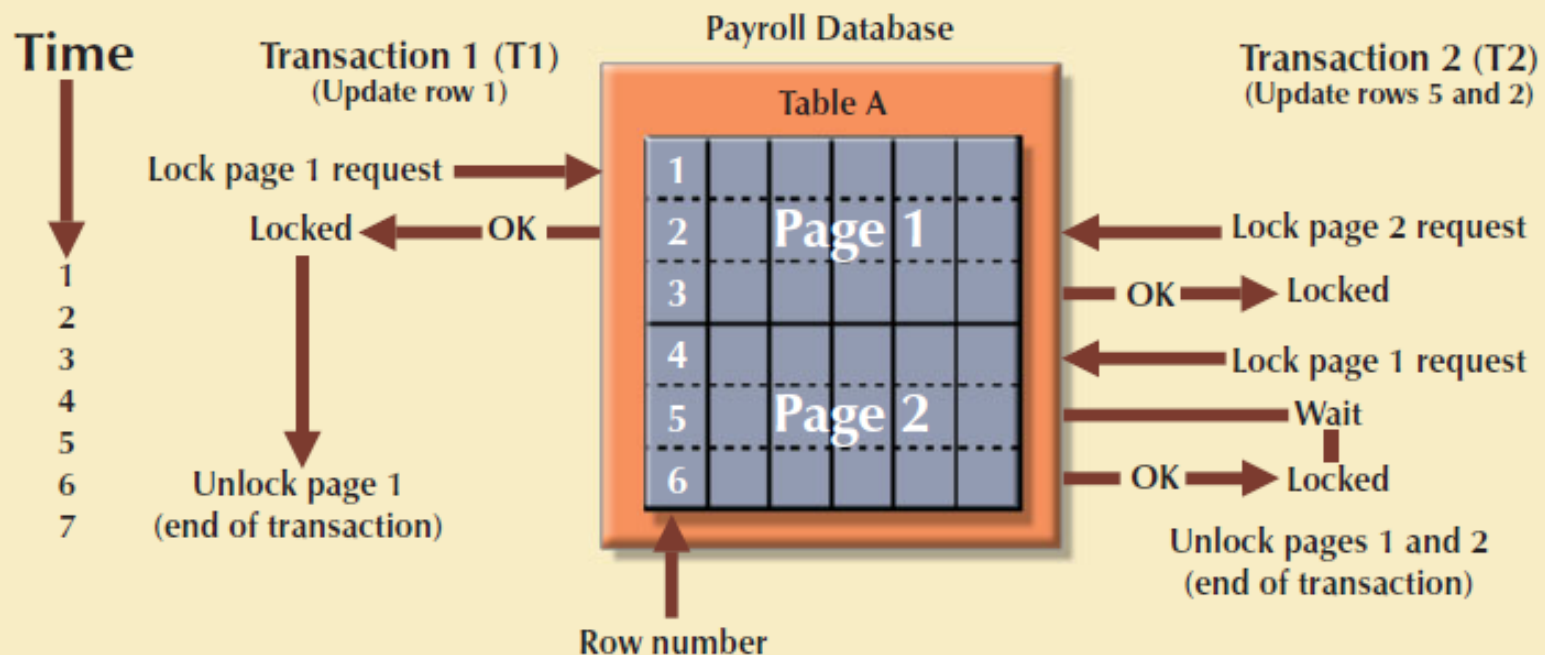
An example of a table-level lock



# Example: Page-level Lock

FIGURE 10.5

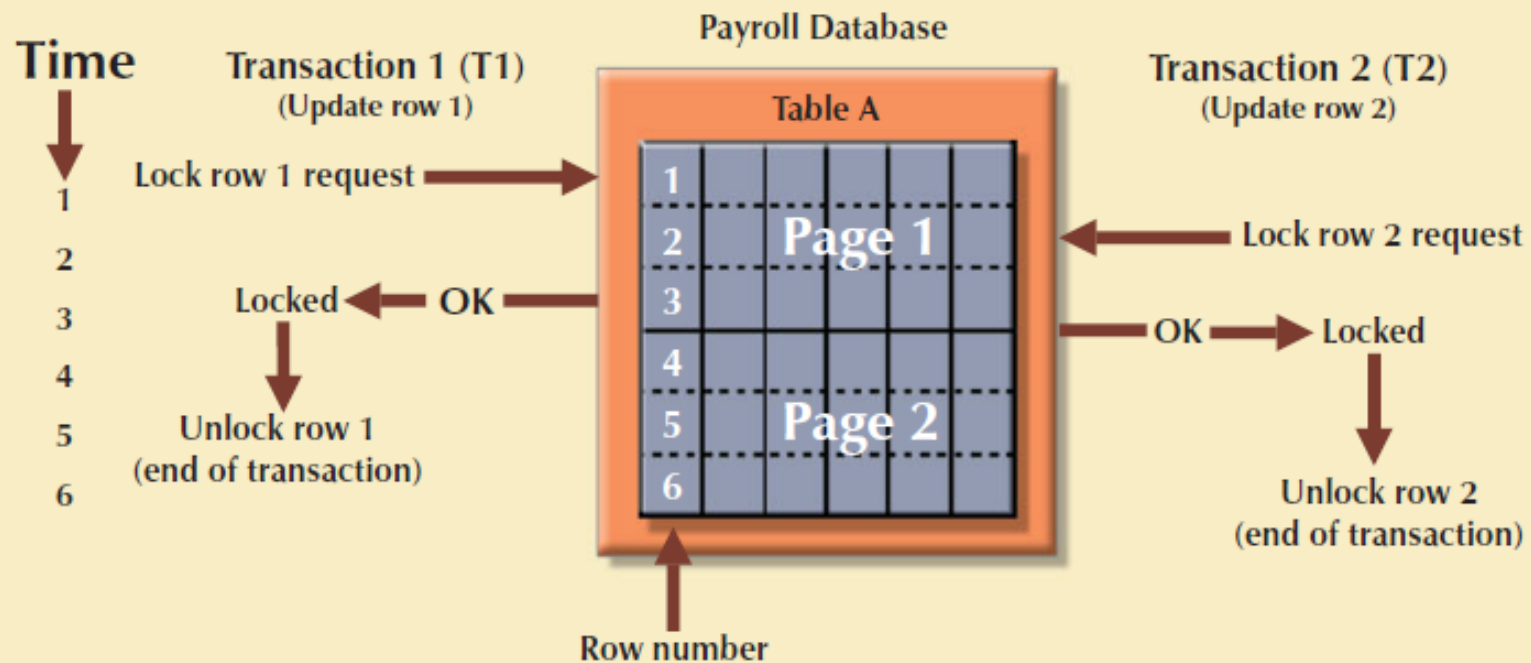
An example of a page-level lock



# Example: Row-level Lock

FIGURE 10.6

An example of a row-level lock



# Lock Types

- [Binary lock](#)
  - Two states: locked (1) or unlocked (0)
  - Every transaction required a lock and unlock operation for each accessed data item, which are automatically managed by the DBMS
- Exclusive lock
  - Access is specifically reserved for transaction that locked object
  - Mutual exclusive rule
  - Must be used when potential for conflict exists
- Shared lock
  - Concurrent transactions are granted read access on basis of a common lock



# Locking Conflict Table

Data Status \ Request	Not Locked	Share Locked	Exclusive Lock
Shared Lock	No Conflict	No Conflict	Conflict
Exclusive Lock	No Conflict	Conflict	Conflict

# Example: Binary Lock

TABLE  
10.12

An Example of a Binary Lock

TIME	TRANSACTION	STEP	STORED VALUE
1	T1	Lock PRODUCT	
2	T1	Read PROD_QOH	13
3	T1	PROD_QOH = $13 + 10$	
4	T1	Write PROD_QOH	23
5	T1	Unlock PRODUCT	
6	T2	Lock PRODUCT	
7	T2	Read PROD_QOH	23
8	T2	PROD_QOH = $23 - 10$	
9	T2	Write PROD_QOH	13
10	T2	Unlock PRODUCT	


# Two-Phase Locking to Ensure Serializability

- Defines how transactions acquire and relinquish locks
- Guarantees serializability, but does not prevent deadlocks
  - Growing phase
    - Transaction acquires all required locks **without unlocking** any data
  - Shrinking phase
    - Transaction releases all locks and **cannot obtain any new lock**
- Governed by the following rules:
  - Two transactions cannot have conflicting locks
  - No unlock operation can precede a lock operation in the same transaction
  - No data are affected until all locks are obtained
    - That is, until transaction in in its locked point

# Two-Phase Locking Protocol

FIGURE 10.7

Two-phase locking protocol



# Deadlocks

- Condition that occurs when two transactions wait for each other to unlock data
- Possible only if one of the transactions wants to obtain an exclusive lock on a data item
  - No deadlock condition can exist among shared locks
- Three techniques to control deadlock:
  - Prevention: Abort & reschedule transaction if deadlock possible
  - Detection: Test database for deadlocks; abort victim, leave rest
  - Avoidance
- Choice of deadlock control method depends on database environment
  - Low probability of deadlock, detection recommended
  - High probability of deadlock, prevention recommended

# How Deadlock Condition is Created?

TABLE  
10.13

How a Deadlock Condition Is Created

TIME	TRANSACTION	REPLY	LOCK STATUS	
0			Data X	Data Y
1	T1:LOCK(X)	OK	Unlocked	Unlocked
2	T2: LOCK(Y)	OK	Locked	Unlocked
3	T1:LOCK(Y)	WAIT	Locked	Locked
4	T2:LOCK(X)	WAIT	Locked	Locked
5	T1:LOCK(Y)	WAIT	Locked	Locked
6	T2:LOCK(X)	WAIT	Locked	Locked
7	T1:LOCK(Y)	WAIT	Locked	Locked
8	T2:LOCK(X)	WAIT	Locked	Locked
9	T1:LOCK(Y)	WAIT	Locked	Locked
...	.....	.....	.....	.....
...	.....	.....	.....	.....
...	.....	.....	.....	.....
...	.....	.....	.....	.....



# Concurrency Control with Time Stamping Methods

- Assign global, unique time stamp to each transaction
- Produces explicit order in which transactions are submitted to DBMS
- Properties of time stamp
  - **Uniqueness**
    - Ensures that no equal time stamp values can exist
  - **Monotonicity**
    - Ensures that time stamp values always increase
- Disadvantages
  - Each value in database needs two fields: a) last time field was read, and b) last update
  - Increases memory needs and processing overhead

# Wait/Die and Wound/Wait Schemes

- Wait/die
  - Older requesting transaction **waits** and
  - Younger requesting transaction is **rolled back** and rescheduled
- Wound/wait
  - Older requesting transaction **preempts (rolls back) younger** transaction and reschedules it
  - Younger requesting transaction **waits**



# Wait/Die and Wound/Wait Schemes (Cont...)

TABLE  
10.14

Wait/Die and Wound/Wait Concurrency Control Schemes

TRANSACTION REQUESTING LOCK	TRANSACTION OWNING LOCK	WAIT/DIE SCHEME	WOUND/WAIT SCHEME
T1 (11548789)	T2 (19562545)	<ul style="list-style-type: none"><li>• T1 waits until T2 is completed and T2 releases its locks.</li></ul>	<ul style="list-style-type: none"><li>• T1 preempts (rolls back) T2.</li><li>• T2 is rescheduled using the same time stamp.</li></ul>
T2 (19562545)	T1 (11548789)	<ul style="list-style-type: none"><li>• T2 dies (rolls back)</li><li>• T2 is rescheduled using the same time stamp</li></ul>	<ul style="list-style-type: none"><li>• T2 waits until T1 is completed and T1 releases its locks.</li></ul>

# Concurrency Control with Optimistic Methods

- Optimistic approach
  - Based on assumption that **majority** of database operations **do not conflict**
  - Does not require locking or time stamping techniques
  - Transaction is **executed without restrictions until it is committed**
  - Phases: read, validation, and write
  - Good for read/query database systems requiring few update transactions; Poor for heavily used DBMS environment

# Reasons for a Crash

- System Crash
- Transaction or System Error
- Local Error or Exception
- Concurrency Control
- Disk Failure
- Catastrophe

# Crash Recovery

- Recovery Manager
  - When DBMS is restarted after crashes, the recovery manager must bring the database to a consistent state
  - Ensures transaction atomicity and durability
  - Undo actions of transactions that do not commit
  - Redo actions of committed transactions during system failures and media failures (corrupted disk).
- Recovery Manager maintains log information during normal execution of transactions for use during crash recovery

# Database Recovery Management

- Restores database to previous consistent state
- Based on atomic transaction property
  - All portions of transaction treated as single logical unit of work
  - All operations applied and completed to produce consistent database
- If transaction operation cannot be completed
  - Transaction aborted
  - Changes to database rolled back (**undone**)

# Concept that Affect Transaction Recovery

- **Write-ahead-log protocol:** ensures transaction logs are written before data is updated
- **Redundant transaction logs:** ensure physical disk failure will not impair ability to recover
- **Buffers:** temporary storage areas in primary memory
- **Checkpoints:** operations in which DBMS writes all its updated buffers to disk

# Transaction Recovery

- Make use of deferred-write and write-through techniques
- **Deferred-write technique**
  - Transaction operations do not immediately update physical database
  - Only transaction log is updated
  - Database is physically updated only after transaction reaches its commit point using transaction log information

# Transaction Recovery (Cont...)

- Recovery process for deferred-write:
  - Identify last checkpoint
  - If transaction **committed before** checkpoint
    - Do nothing
  - If transaction committed after checkpoint
    - Use transaction log **to redo** the transaction
  - If transaction had ROLLBACK operation or was left active
    - **Do nothing because no updates were made**



# Transaction Recovery (Cont...)

- **Write-through technique**
  - Database is **immediately updated** by transaction operations during transaction's execution
- Recovery process for write-through
  - Identify last checkpoint
  - If transaction committed before checkpoint
    - Do nothing
  - If transaction committed after checkpoint
    - DBMS redoes the transaction using "after" values
  - If transaction had ROLLBACK operation or was left active
    - **Uses the before value in the transaction log records to ROLLBACK (undo)**

# Transaction Recovery (Cont...)

TABLE  
10.15

A Transaction Log for Transaction Recovery Examples

TRL ID	TRX NUM	PREV PTR	NEXT PTR	OPERATION	TABLE	ROW ID	ATTRIBUTE	BEFORE VALUE	AFTER VALUE
341	101	Null	352	START	****Start Transaction				
352	101	341	363	UPDATE	PRODUCT	54778-2T	PROD_QOH	45	43
363	101	352	365	UPDATE	CUSTOMER	10011	CUST_BALANCE	615.73	675.62
365	101	363	Null	COMMIT	**** End of Transaction				
397	106	Null	405	START	****Start Transaction				
405	106	397	415	INSERT	INVOICE	1009			1009,10016, ...
415	106	405	419	INSERT	LINE	1009,1			1009,1, 89-WRE-Q,1, ...
419	106	415	427	UPDATE	PRODUCT	89-WRE-Q	PROD_QOH	12	11
423				CHECKPOINT					
427	106	419	431	UPDATE	CUSTOMER	10016	CUST_BALANCE	0.00	277.55
431	106	427	457	INSERT	ACCT_TRANSACTION	10007			1007,18-JAN-2008, ...
457	106	431	Null	COMMIT	**** End of Transaction				
521	155	Null	525	START	****Start Transaction				
525	155	521	528	UPDATE	PRODUCT	2232/QWE	PROD_QOH	6	26
528	155	525	Null	COMMIT	**** End of Transaction				
***** C *R*A* S* H *****									

# Summary

- Transaction: sequence of database operations that access database
  - Logical unit of work
    - No portion of transaction can exit by itself
  - Five main properties: atomicity, consistency, isolation, durability, and serializability
- COMMIT saves changes to disk
- ROLLBACK restores previous database state
- SQL transactions are formed by several SQL statements or database requests

# Summary (Cont...)

- Transaction log keeps track of all transactions that modify database
- Concurrency control coordinates simultaneous execution of transactions
- Scheduler establishes order in which concurrent transaction operations are executed
- Lock guarantee unique access to data item by transaction
- Two types of locks: binary locks and shared/exclusive locks

# Summary (Cont...)

- Serializability of schedules is guaranteed through the use of two-phase locking
- Deadlock: when two or more transactions wait indefinitely for each other to release lock
- Three deadlock control techniques: prevention, detection, and avoidance
- Time stamping methods assign unique time stamp to each transaction
  - Schedules execution of conflicting transactions in time stamp order

# Summary (Cont...)

- Optimistic methods assume the majority of database transactions do not conflict
  - Transactions are executed concurrently, using private copies of the data
- Database recovery restores database from given state to previous consistent state